University of Glasgow

Diffusion models for Image Restoration

Daniela Ivanova

Computer Vision and Autonomous Systems (CVAS)

University of Glasgow

What is image restoration?

Colourisation

Super-resolution

Inpainting (regular shape)

Inpainting (arbitrary artefacts)

Diffusion architecture - recap

- Since the *input, output and latents have the same dimensionality*, for each step the denoising model is a U-Net with some modifications
  - Group normalisation
  - Global self-attention
  - Sinusoidal positional **time embeddings** concatenated to the input of each block
- Instead of predicting the denoised image, the network predicts the noise that was added to it

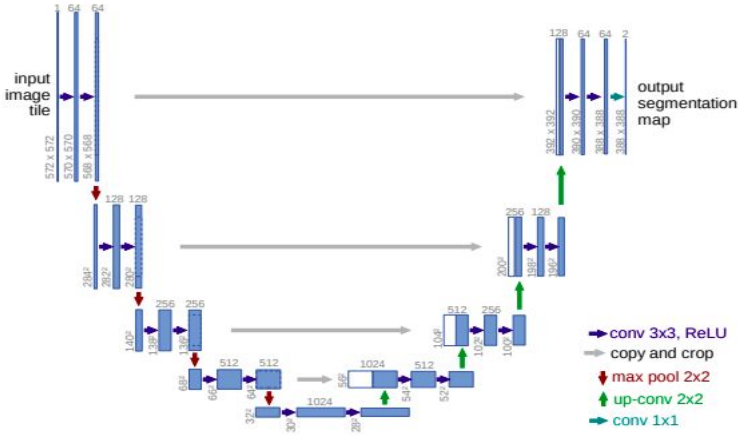

Credit: Ronneberger et al, “Convolutional Networks for Biomedical Image Segmentation”

University of Glasgow

Diffusion process - recap

Distribution of noised images

$$q(x_1, ..., x_T|x_0) := \prod_{t=1}^{T} q(x_t|x_{t-1}) \quad (1)$$

scaling

$$q(x_t|x_{t-1}) := \mathcal{N}(x_t; \sqrt{1-\beta_t}x_{t-1}, \beta_t\mathbf{I}) \quad (2)$$

diagonal covariance matrix

- To produce each latent, we can add noise iteratively (slow)

University of Glasgow

Diffusion process - recap

noised latents directly conditioned on the input x_0. With $\alpha_t := 1 - \beta_t$ and $\bar{\alpha}_t := \prod_{s=0}^{t} \alpha_s$, we can write the marginal

$$q(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)\mathbf{I}) \quad (8)$$

$$x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon \quad (9)$$

where $\epsilon \sim \mathcal{N}(0, \mathbf{I})$. Here, $1 - \bar{\alpha}_t$ tells us the variance of the noise for an arbitrary timestep, and we could equivalently use this to define the noise schedule instead of β_t.

- …or we can derive the variance scale for an arbitrary step - accumulate the noise from the first step to the step we need

University of Glasgow

- We can condition on class labels - embed class label v_i along with time embedding e_t - (Nichol & Dhariwal, 2021)

Credit: Nichol & Dhariwal, "Improved denoising diffusion probabilistic models."

University of Glasgow

Diffusion conditioning - recap

- Classifier guidance (Dhariwal & Nichol, 2021)

$$p_{\theta,\phi}(x_t|x_{t+1},y) = Z p_\theta(x_t|x_{t+1}) p_\phi(y|x_t)$$

$$\begin{aligned}\nabla_{x_t} \log(p_\theta(x_t) p_\phi(y|x_t)) &= \nabla_{x_t} \log p_\theta(x_t) + \nabla_{x_t} \log p_\phi(y|x_t) \\ &= -\frac{1}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(x_t) + \nabla_{x_t} \log p_\phi(y|x_t)\end{aligned}$$

$$\hat{\epsilon}(x_t) := \epsilon_\theta(x_t) - \sqrt{1-\bar{\alpha}_t}\, \nabla_{x_t} \log p_\phi(y|x_t)$$

University of Glasgow

Diffusion conditioning - recap

- Classifier guidance (Dhariwal & Nichol, 2021)
  - y comes from the downsampling half of the UNet, which is used as a classifier

Figure 6: Samples from BigGAN-deep with truncation 1.0 (FID 6.95, left) vs samples from our diffusion model with guidance (FID 4.59, middle) and samples from the training set (right).

Credit: Dhariwal & Nichol, "Diffusion Models Beat GANs on Image Synthesis"

University of Glasgow

Diffusion conditioning - recap

"a hedgehog using a calculator"

"a corgi wearing a red bowtie and a purple party hat"

"robots meditating in a vipassana retreat"

"a fall landscape with a small cottage next to a lake"

"a surrealist dream-like oil painting by salvador dalí of a cat playing checkers"

"a professional photo of a sunset behind the grand canyon"

"a high-quality oil painting of a psychedelic hamster dragon"

"an illustration of albert einstein wearing a superhero costume"

- We can condition on text descriptions
  - Each attention layer is attending to each token for the text embedding
  - Doesn't work very well still

Credit: Dhariwal & Nichol, "GLIDE: Towards Photorealistic Image Generation and Editing with Text-Guided Diffusion Models"

University of Glasgow

Diffusion conditioning - recap

- CLIP guided diffusion:
  - At inference time, use CLIP guidance
  - CLIP outputs a similarity score between image and text for each pixel
  - Use that gradient at each time step to push the image in the direction which would give it higher score/smaller CLIP loss

(c) GLIDE (CLIP guidance, scale 2.0)

University of Glasgow

Diffusion conditioning - recap

- Classifier-free guidance:
  - Train with and without text embeddings
  - Predict an image without the text prompt and with the text prompt
  - Find the difference between the two
  - Use that gradient to go in the direction of the image with text using a scaling factor for the vector

(d) GLIDE (Classifier-free guidance, scale 3.0)

University of Glasgow

Diffusion models for image restoration

End-to-end training with conditioning

- SR3
- Palette

Using pre-trained models, conditioning only during inference

- DDRM
- RePaint
- Stable Diffusion
- DiffEdit (bonus!)

University of Glasgow

SR3 (Saharia et al.)

- We can condition on another image
  - low res image for superresolution
  - grayscale image for colourisation
  - images with missing patches for inpainting
- Concatenate noise vector with the conditioning image
- Slow, diffusing the entire image

Credit: Saharia et al, “Image Super-Resolution via Iterative Refinement”

University of Glasgow

alette: Image-to-Image Diffusion Models (Saharia et al.)

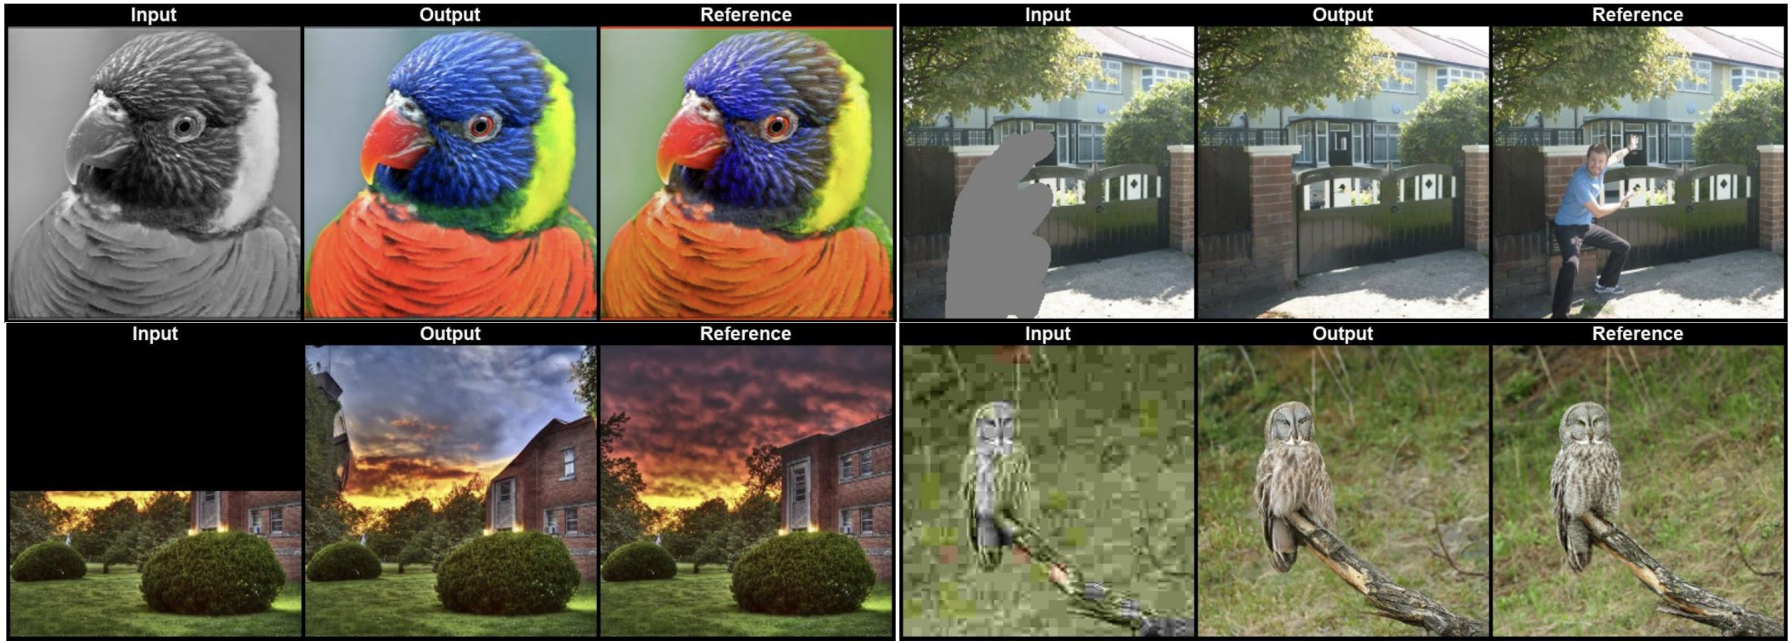

Credit: Saharia et al, “Palette: Image-to-Image Diffusion Models”

Diffusion model for film artifact removal

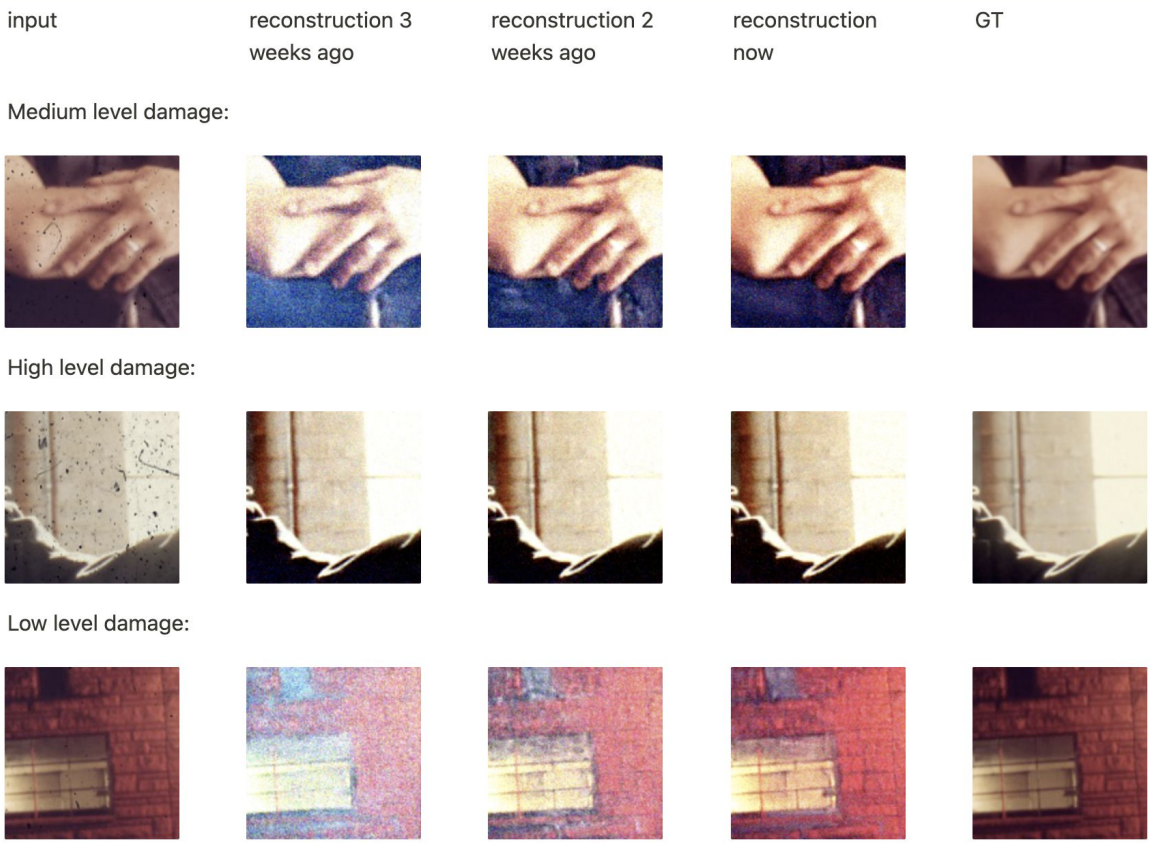

- 97.8 M params
- bs=8
- 1M iterations in SR3 paper
- ***very*** slow if you don't have a TPU
- task-specific - need lots of data

University of Glasgow

DDRM (Kawar et al., 2022)

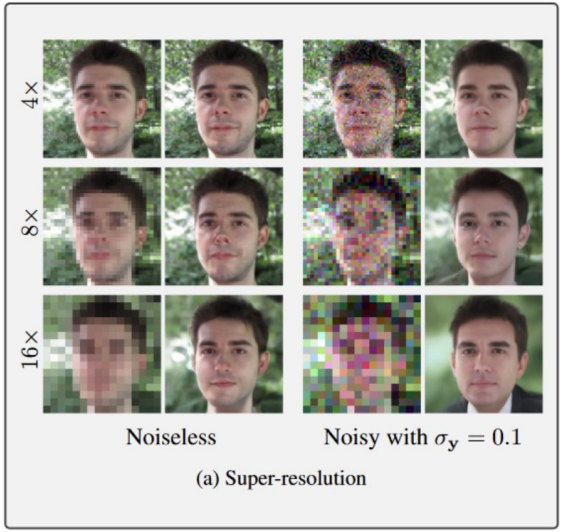

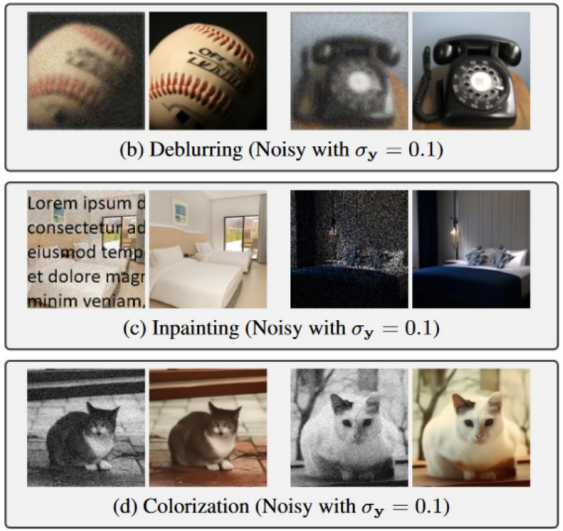

- use pretrained unconditional DDPM
- decompose degradation operator ***H*** using SVD
- perform diffusion in spectral space

Credit: Kawar et al, “Denoising Diffusion Restoration Models”

$$\mathbf{y} = \boldsymbol{H}\mathbf{x} + \mathbf{z}, \qquad q(\mathbf{x}_{1:T}|\mathbf{x}_0, \mathbf{y}) = q^{(T)}(\mathbf{x}_T|\mathbf{x}_0, \mathbf{y}) \prod_{t=0}^{T-1} q^{(t)}(\mathbf{x}_t|\mathbf{x}_{t+1}, \mathbf{x}_0, \mathbf{y}),$$

RePaint (Lugmayr et al., 2022)

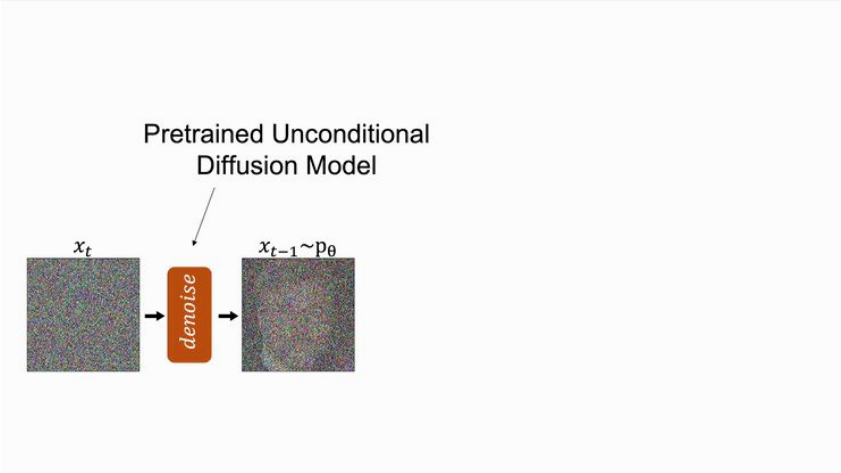

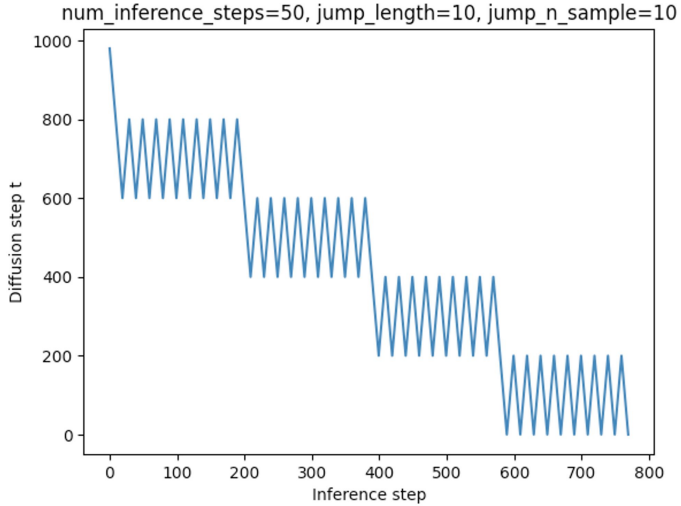

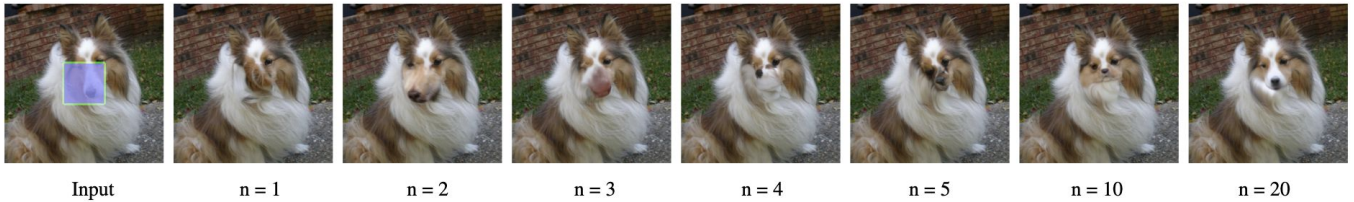

Credit: Lugmayr et al, "RePaint: Inpainting using Denoising Diffusion Probabilistic Models"

Figure 3. **The effect of applying n sampling steps.** The first example with $n = 1$ is the DDPM baseline, the second with $n = 2$ is with one resample step. More resampling steps lead to more harmonized images. The benefit saturates at about $n = 10$ resamplings.

University of Glasgow

Latent diffusion (Romach et al.)

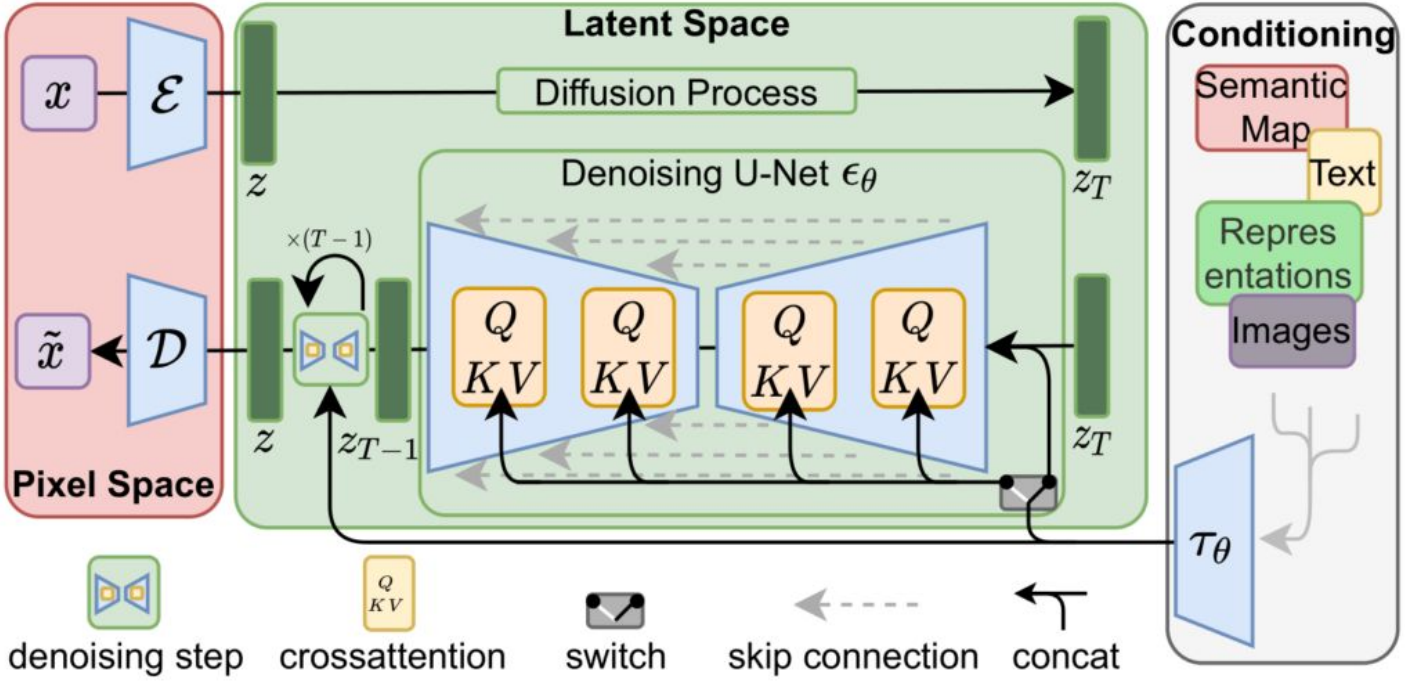

Credit: Romach et al., “High-Resolution Image Synthesis with Latent Diffusion Models”

University of Glasgow

Latent diffusion (Romach et al.)

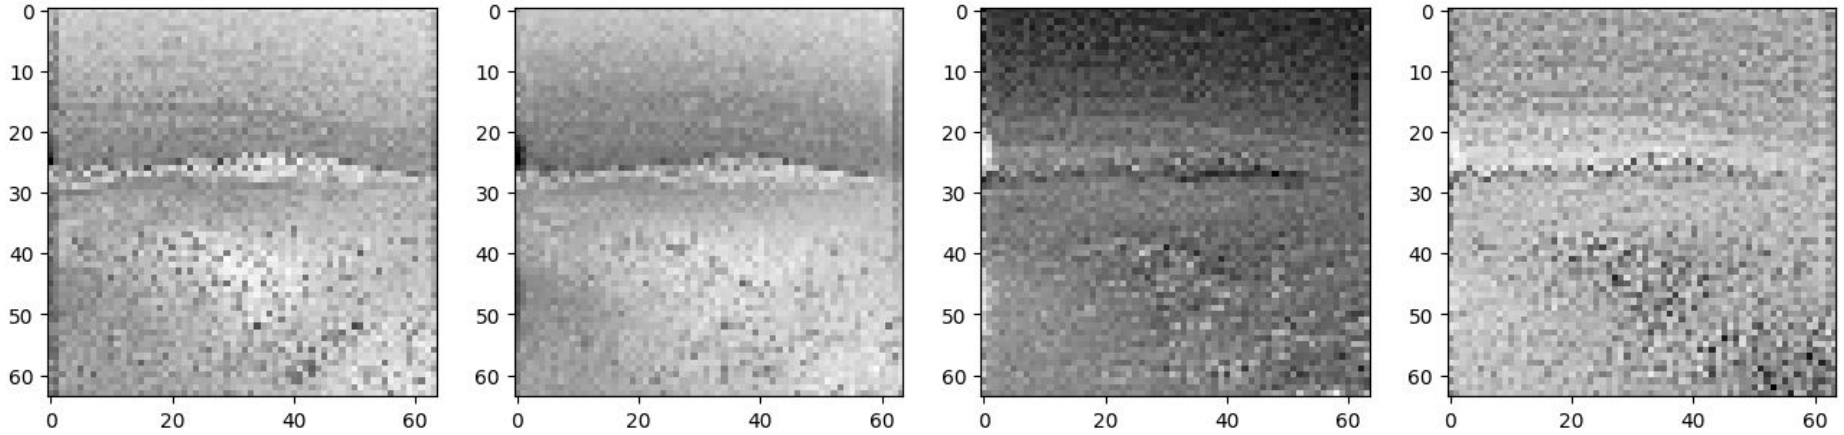

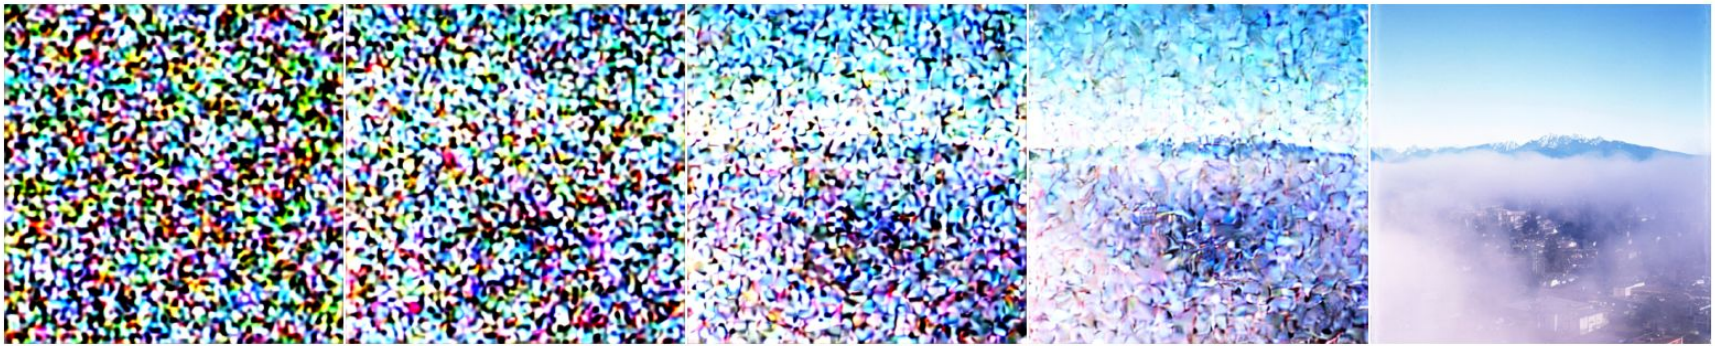

University of Glasgow

Stable Diffusion Inpainting

“cat sitting on a bench”

University of Glasgow

Stable Diffusion Inpainting

Unconditional (legacy)

- use pre-trained SD
- make prediction from noise
- mask out the latents
- make next prediction
- etc

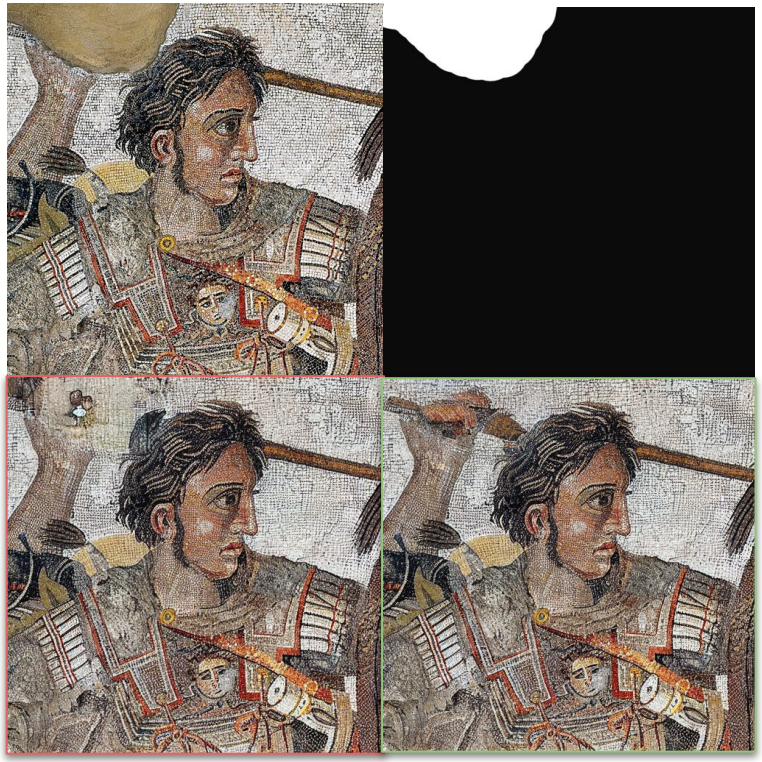

Conditional

- fine-tune pre-trained SD model on inpainting
- pass the mask, masked latents, original latents as a 9-channel input to the U-Net

University of Glasgow

RePaint (Lugmayr et al., 2022)

What if we tried RePaint it in latent space?

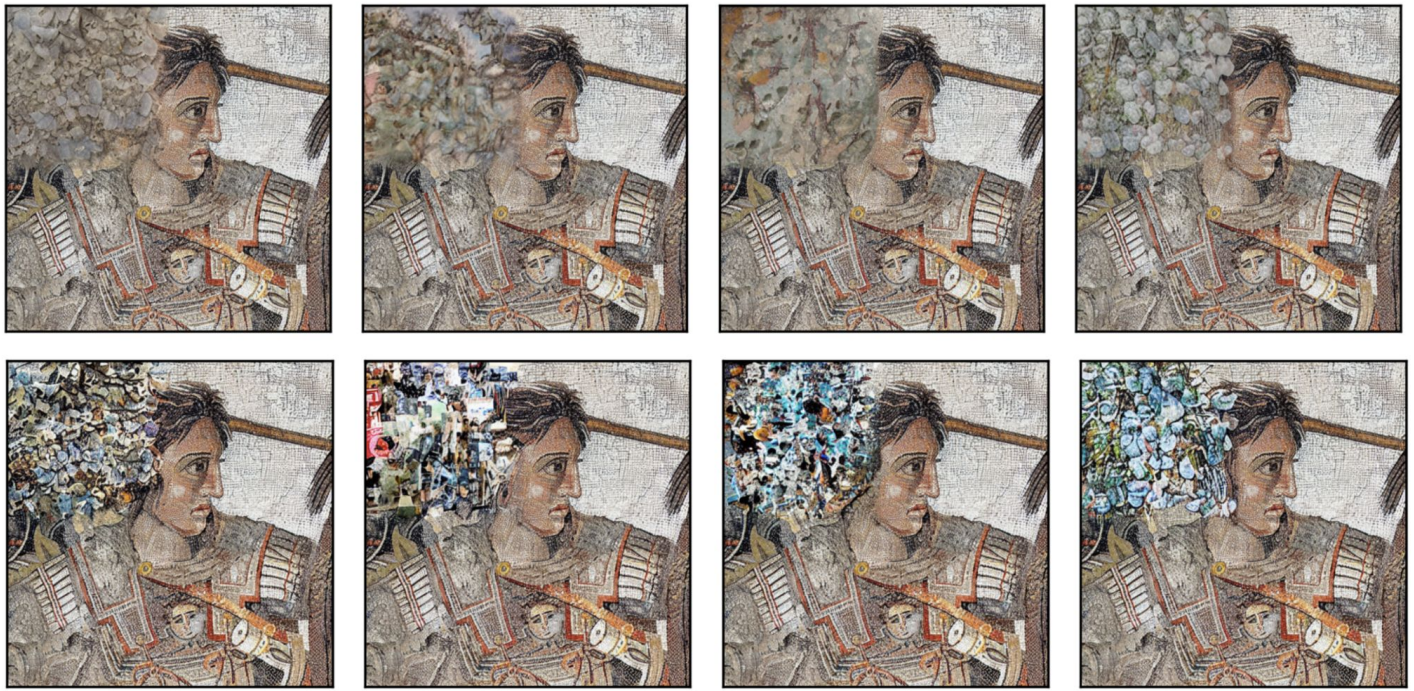

University of Glasgow

RePaint (Lugmayr et al., 2022)

What if we tried it in latent space?

- Only works with the sampler used to train SD

DiffEdit (Couairon et al., 2022)

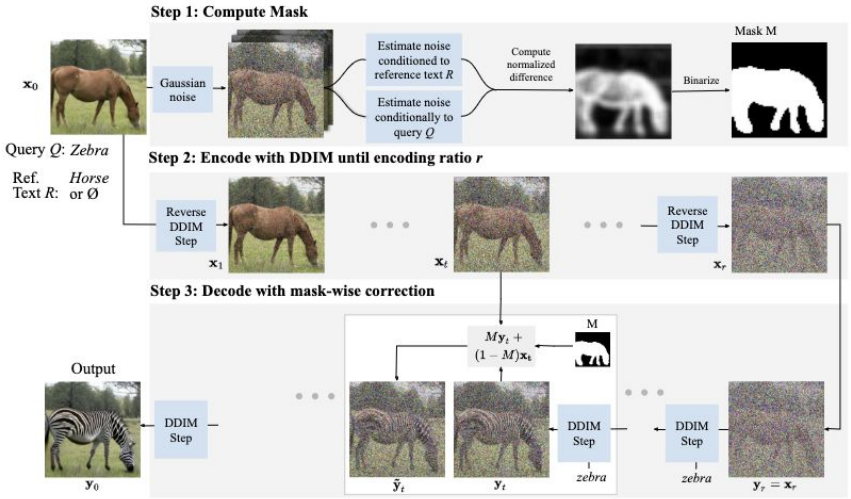

Credit: Couairon et al.,"DiffEdit: Diffusion-based semantic image editing with mask guidance"

- denoise once using reference text
- denoise again using query text
- the difference in noise estimates => locations that are predicted to change the most between conditioning on the original and new texts

University of Glasgow

DiffEdit (Couairon et al., 2022)

reference: “horse”
query: “zebra”

- noise-denoise 10 times with each prompt
- accumulate predicted noises
- find difference

Check out “DiffEdit paper implementation” by Kevin Bird

DiffEdit (Couairon et al., 2022)

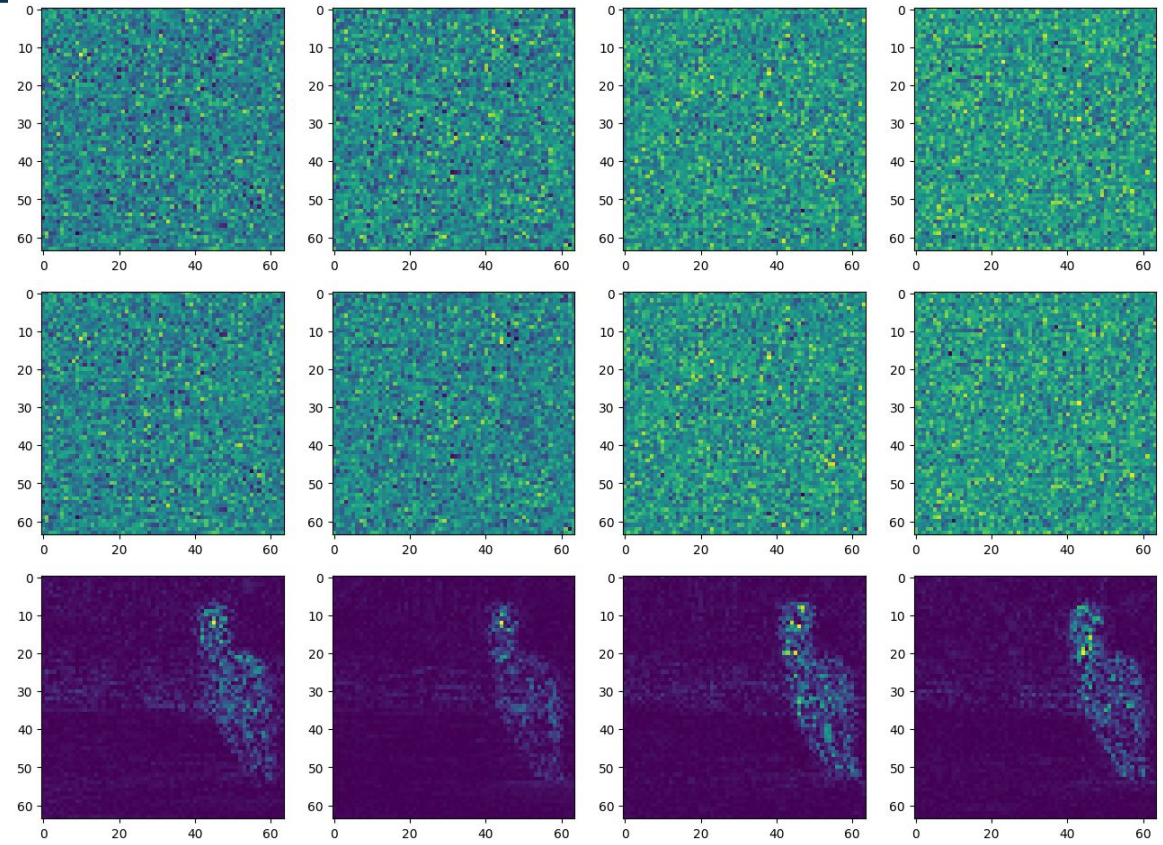

“horse” noise

“zebra” noise

difference

DiffEdit (Couairon et al., 2022) - new idea

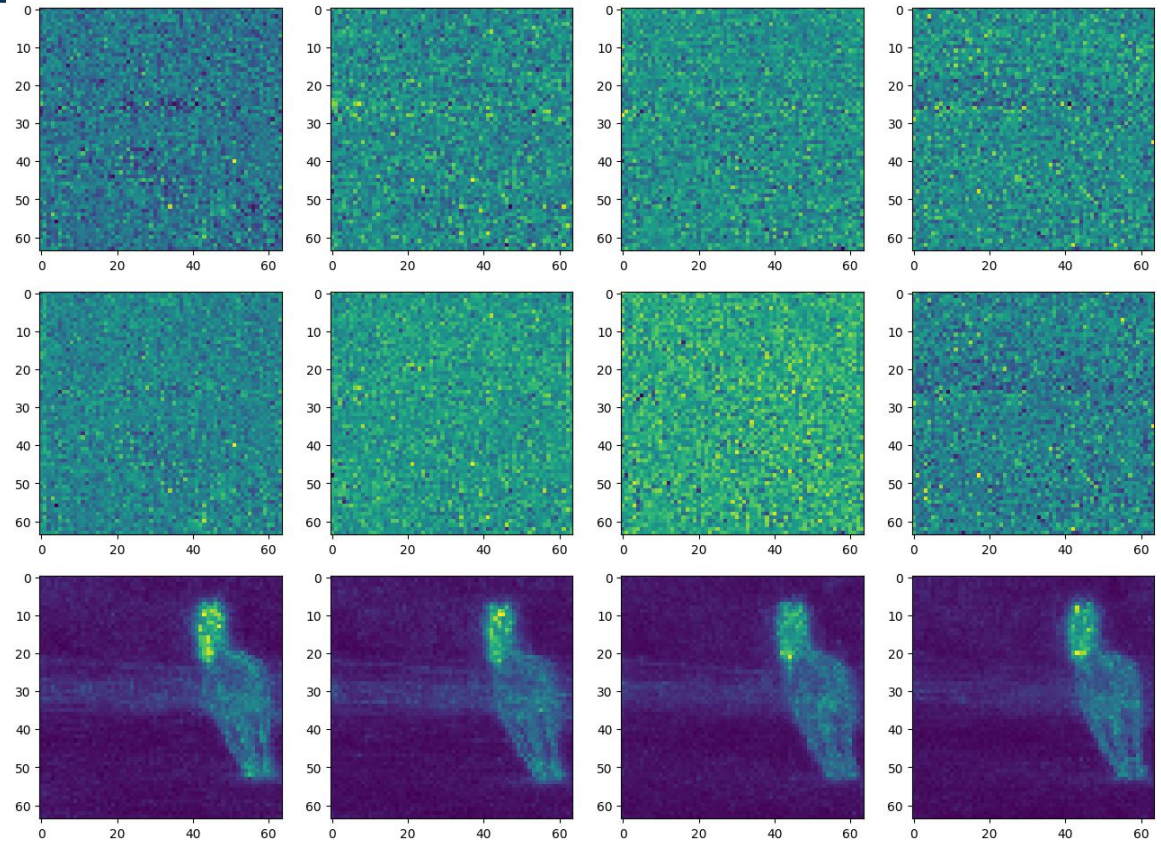

“horse” noise

“- horse” noise

difference

University of Glasgow

DiffEdit (Couairon et al., 2022)

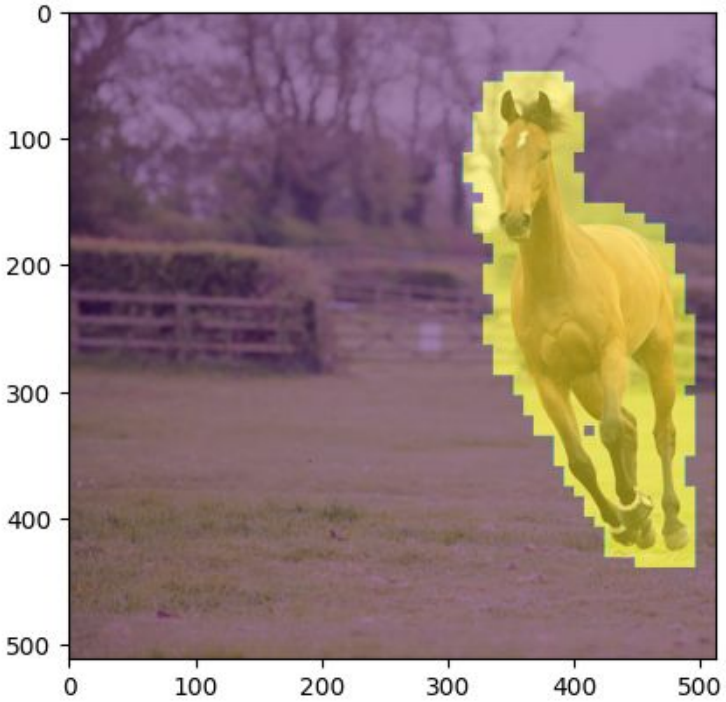

horse - zebra

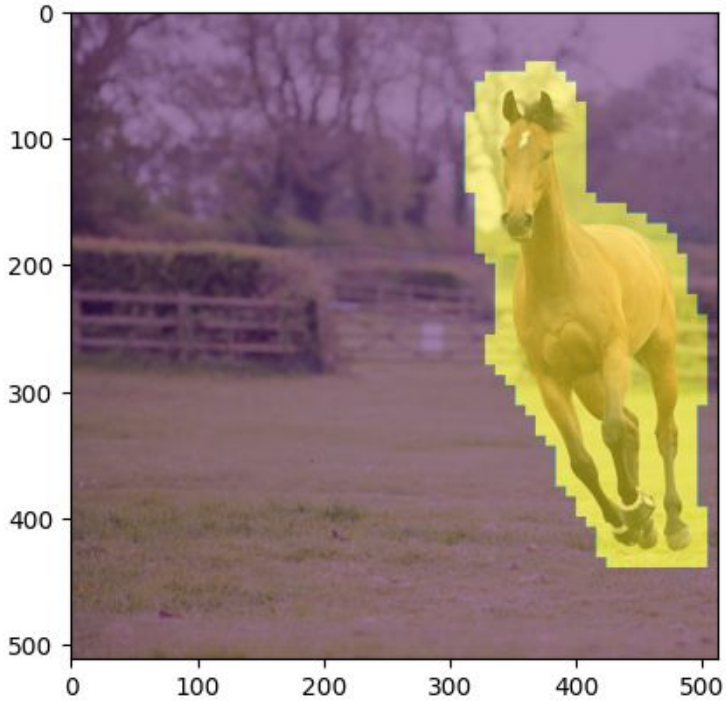

horse - (-horse)

University of Glasgow

DiffEdit (Couairon et al., 2022)

damaged film photo

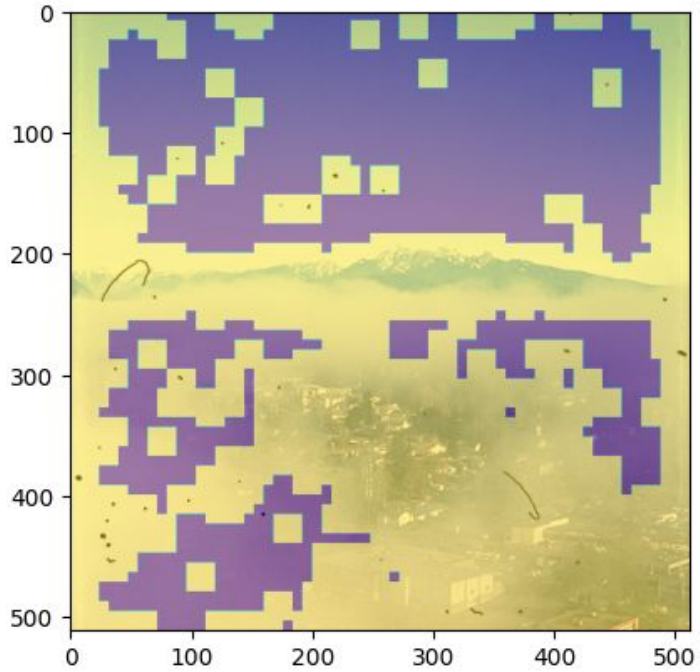

damaged film photo - (-damaged film photo)

University of Glasgow

Thank you!

University of Glasgow

References

Ronneberger, Olaf, Philipp Fischer, and Thomas Brox. "U-net: Convolutional networks for biomedical image segmentation." *Medical Image Computing and Computer-Assisted Intervention–MICCAI 2015: 18th International Conference, Munich, Germany, October 5-9, 2015, Proceedings, Part III 18*. Springer International Publishing, 2015.

Ho, Jonathan, Ajay Jain, and Pieter Abbeel. "Denoising diffusion probabilistic models." *Advances in Neural Information Processing Systems* 33 (2020): 6840-6851.

Song, Jiaming, Chenlin Meng, and Stefano Ermon. "Denoising diffusion implicit models." *arXiv preprint arXiv:2010.02502* (2020).

Nichol, Alexander Quinn, and Prafulla Dhariwal. "Improved denoising diffusion probabilistic models." *International Conference on Machine Learning*. PMLR, 2021.

Dhariwal, Prafulla, and Alexander Nichol. "Diffusion models beat gans on image synthesis." *Advances in Neural Information Processing Systems* 34 (2021): 8780-8794.

Nichol, Alex, et al. "Glide: Towards photorealistic image generation and editing with text-guided diffusion models." *arXiv preprint arXiv:2112.10741* (2021).

Saharia, Chitwan, et al. "Image super-resolution via iterative refinement." *IEEE Transactions on Pattern Analysis and Machine Intelligence* (2022).

Saharia, Chitwan, et al. "Palette: Image-to-image diffusion models." *ACM SIGGRAPH 2022 Conference Proceedings*. 2022.

Kawar, Bahjat, et al. "Denoising diffusion restoration models." *arXiv preprint arXiv:2201.11793* (2022).

Lugmayr, Andreas, et al. "Repaint: Inpainting using denoising diffusion probabilistic models." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2022.

Rombach, Robin, et al. "High-resolution image synthesis with latent diffusion models." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2022.

Couairon, Guillaume, et al. "Diffedit: Diffusion-based semantic image editing with mask guidance." *arXiv preprint arXiv:2210.11427* (2022).